

# CS420: Operating Systems

## Virtual Memory

---

James Moscola

Department of Engineering & Computer Science

York College of Pennsylvania



# Background

---


- **Code needs to be in memory to execute, but entire program rarely used**
  - Error code, unusual routines, large data structures
- **Want the ability to execute partially-loaded program**
  - Programs no longer constrained by limits of physical memory
  - Program data no longer constrained by limits of physical memory

# Background - Virtual Memory

---


- **Virtual memory** – separation of user logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows physical address spaces to be shared by several processes
  - Allows for more efficient process creation
  - More programs running concurrently
  - Less I/O needed to load or swap processes
- In contrast to **dynamic loading**, virtual memory does not require programmer to do anything extra

# Virtual Memory that is Larger than Physical Memory



# Virtual Address Space

- Enables **sparse address spaces with holes left for growth, dynamically linked libraries, etc.**
  - e.g. don't waste physical memory with empty space that is intended for the growth of the stack/heap
- **System libraries can be shared by mapping them into virtual address space**
- **Can create shared memory by mapping pages into virtual address space**
- **Virtual memory allows pages to be shared during fork(), speeding up process creation**



# Demand Paging

---

- Could bring entire process into memory at load time
- Or load a page into memory only when it is needed, called demand paging
  - Pages that are never used are never loaded
    - Less I/O needed, no unnecessary I/O
    - Less memory needed
    - Faster response
  - Similar to page table system with swapping, but demand paging doesn't swap entire process into memory, instead uses a lazy swapper
- Lazy swapper – never swaps a page into memory unless it will be needed
  - A swapper that deals with pages is called a pager

# Page Table Example

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical  
memory

valid-invalid bit	
frame	bit
0	4 v
1	i
2	6 v
3	i
4	i
5	9 v
6	i
7	i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory



Not all pages are always  
resident in physical memory

# Valid-Invalid Bit

---

- **Each page table entry includes a valid-invalid bit**
  - If valid, then process is allowed to access that page, AND it is in physical memory
  - If invalid, process may not be allowed to access the requested page, or the page may not yet be in physical memory

- **Initially valid-invalid bit is set to i on all entries**

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

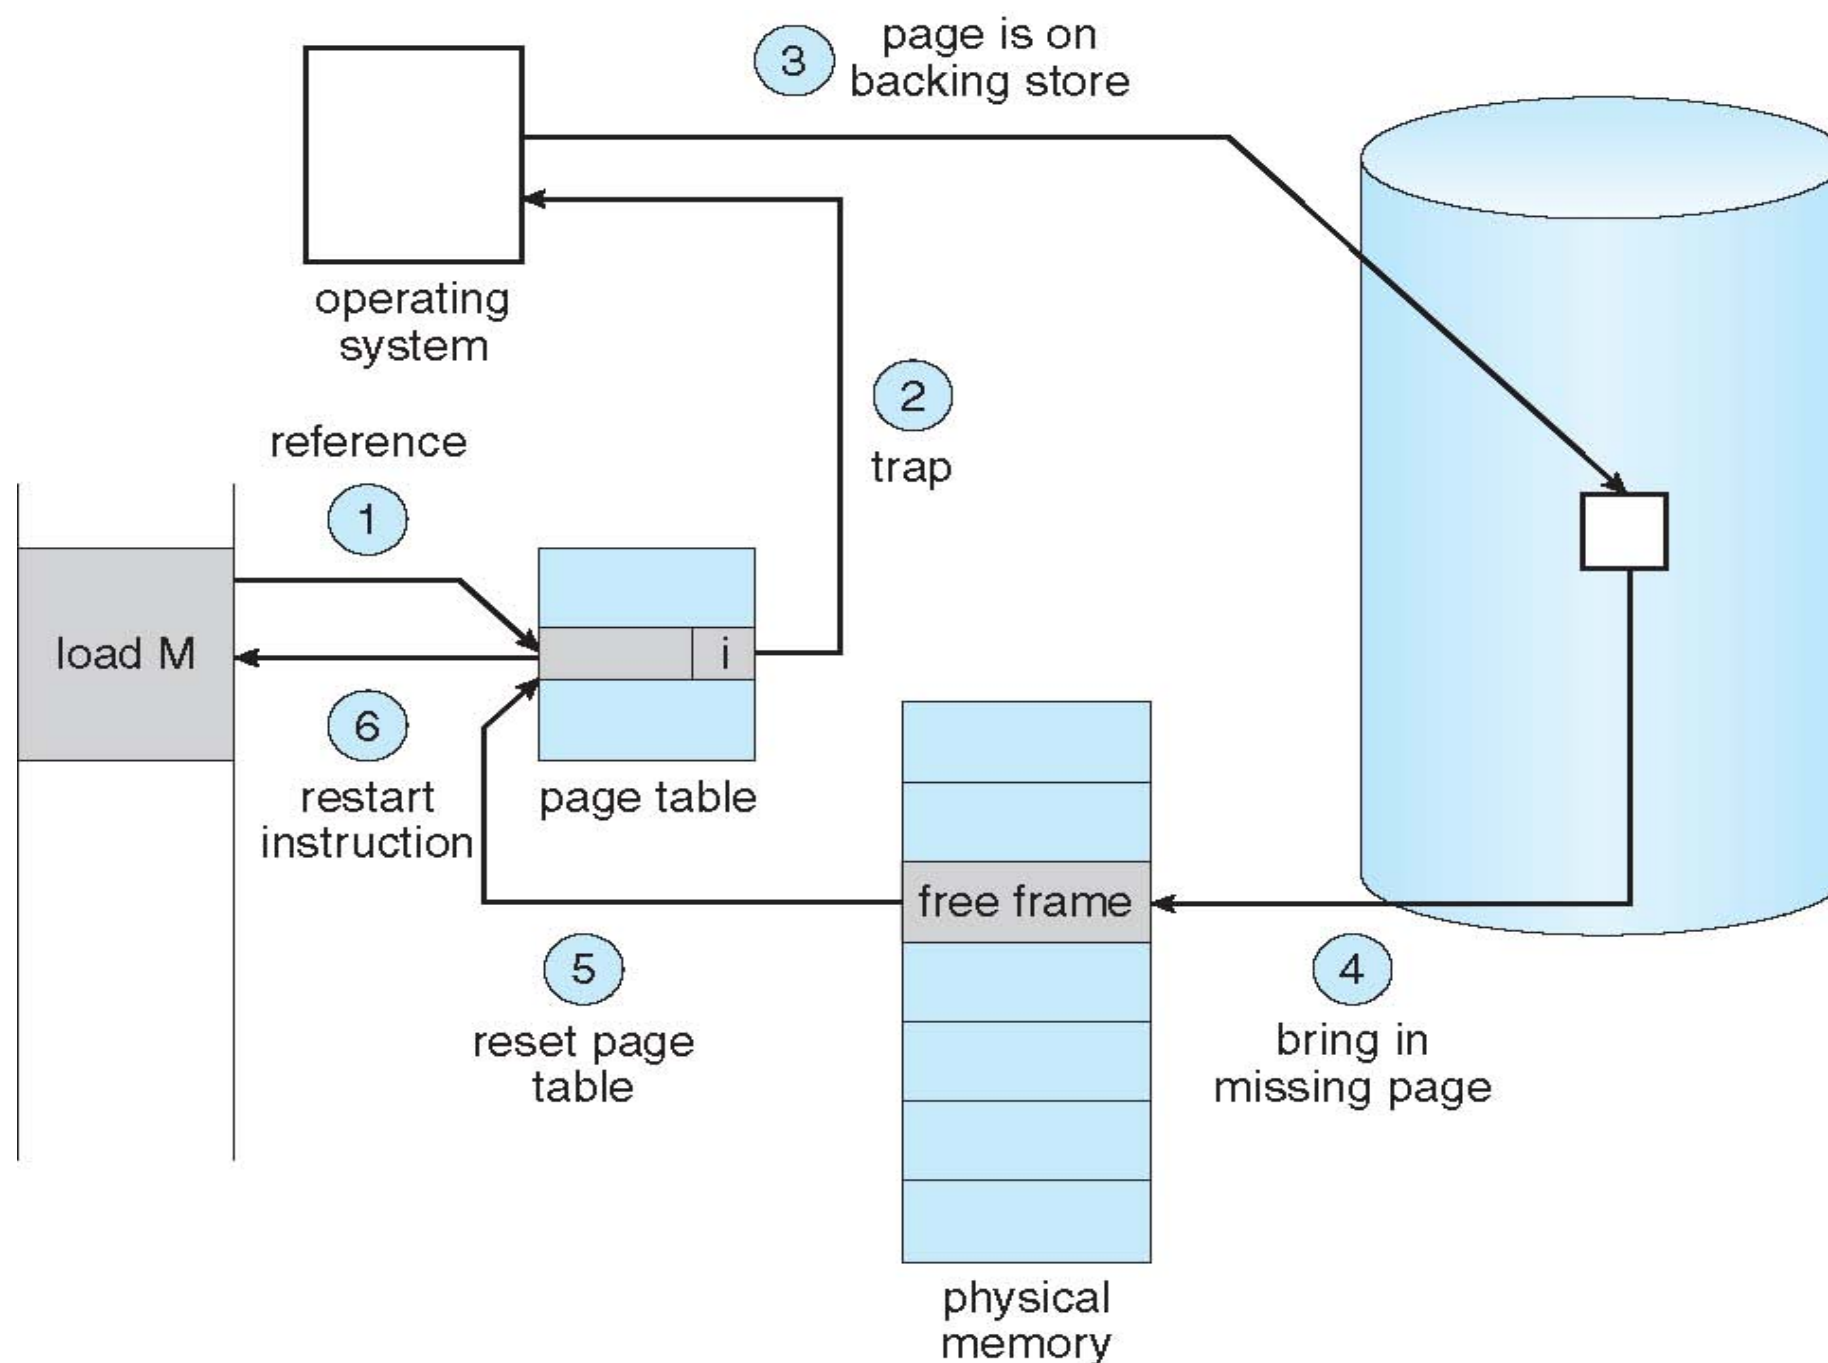


# Page Fault

---

- If requested page is **memory resident**, then process can operate normally
- If requested page not in physical memory (i.e. page is invalid in page table), then a **page fault** occurs
- If page fault occurs, first check to see if the requested page was an illegal request or if it just isn't in physical memory
  - If an illegal reference, the process terminates (seg fault, bus error, ...)
  - If simply not in memory:
    - Get a free frame of physical memory from the free-list
    - Swap page from disk into the frame via a scheduled disk operation
    - Update tables to indicate that the page is now in physical memory (i.e. set valid-invalid bit to 'v')
    - Restart the instruction that caused the page fault

# Steps in Handling a Page Fault



# Aspects of Demand Paging

---

- **Extreme case of demand paging – start a process with no pages in memory**
  - Never bring a page into physical memory until it is needed
  - A page fault will occur each time a new page is requested, including on the very first page
  - The scheme is called **pure demand paging**
- **Demand paging (and pure demand paging) require hardware support**
  - Page table with valid / invalid bit
  - Secondary memory (swap device with swap space)
  - Ability to restart an instruction
- **Performance of system can be severely impacted if too much page swapping is required (see example in OSC9 section 9.2.2, 40x slowdown!)**


# Page Sharing - Copy-on-Write

---

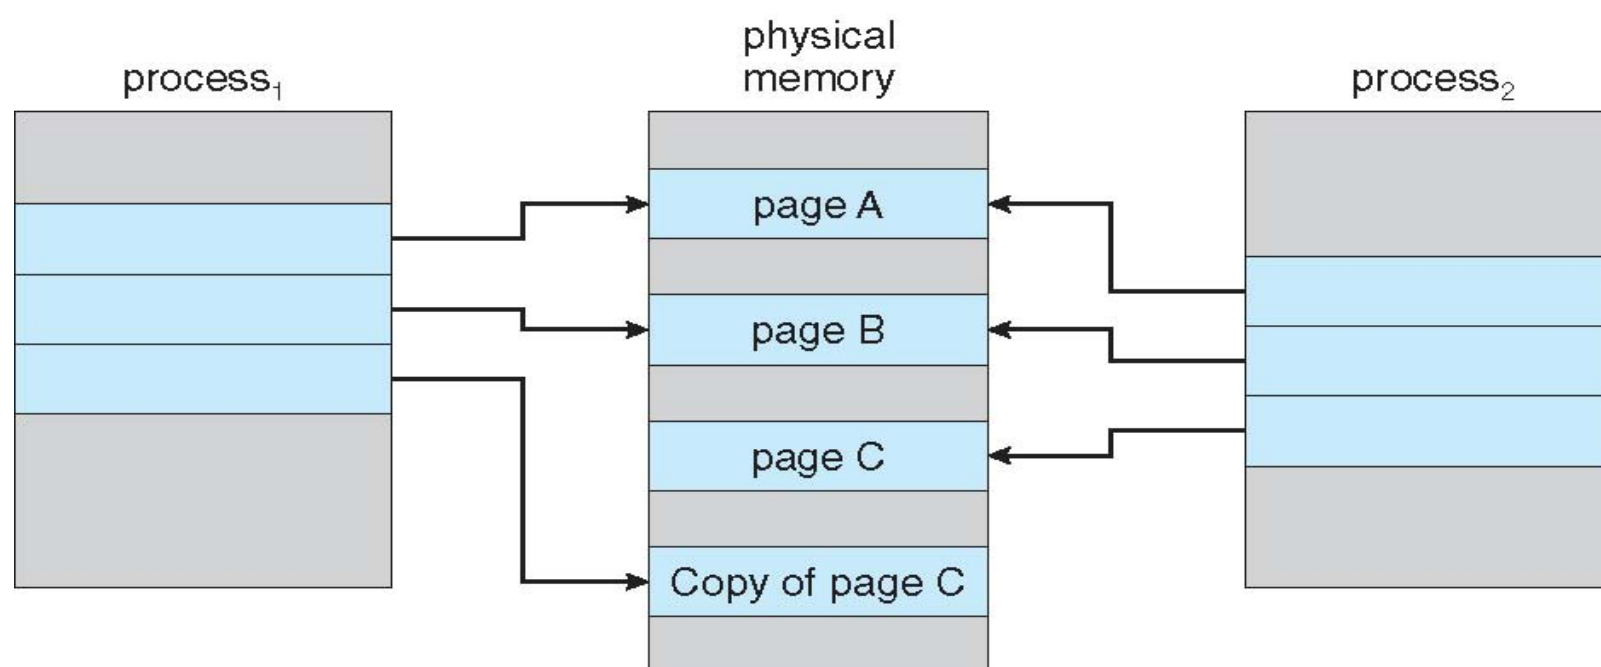
- During process creation **Copy-On-Write (COW)** allows both parent and child processes to initially share the same pages in memory
  - If either process modifies a shared page, only then is the page copied
- **COW allows more efficient process creation as only modified pages are copied**
- **Free pages are typically allocated from a pool of zero-fill-on-demand pages**
- **A variation of `fork()` exists called `vfork()` (virtual memory fork)**
  - Parent suspends so child can use its memory address space
  - If child modifies parent's address space, those changes will be visible to parent (i.e. it does NOT use copy-on-write)
  - Designed to be used when child calls `exec()` immediately after `fork()`
  - Very efficient since no pages are copied during creation of child

# Example of Copy-on-Write

Memory before either process tries to write to shared pages



Memory after process<sub>1</sub> writes to page C



# What Happens if There is no Free Frame?

---


- **All frames used up by process pages and I/O buffers**
  - How much memory should be allocated to I/O and how much to processes?
- **Page replacement – find some page in physical memory, but not really in use, page it out to swap space**
  - Once a page is in physical memory, that doesn't mean it will always be in physical memory
  - Same page may be brought into physical memory several times
  - For performance reasons, want a page replacement algorithm which will result in minimum number of page faults

# Page Replacement

---

- **Must prevent over-allocation of physical memory by modifying page-fault service routine to include page replacement**
  - If no free frame exists, must select a victim frame and write the page that it contains off to swap space
  - After writing page to swap space, must update pages tables accordingly
  - Required page can then be read into newly freed frame
  - Continue process by restarting instruction that caused the page-fault
- **Use modify (dirty) bit to reduce overhead of page transfers**
  - Only modified pages are written back to disk from physical memory
  - Without modify bit, unmodified pages may be unnecessarily written back to swap space

# Need For Page Replacement





# Page and Frame Replacement Algorithms

---

- **Frame-allocation algorithm determines**
  - How many frames to give each process
- **Page-replacement algorithm determines**
  - Which frames to replace
  - Want lowest page-fault rate on both first access and re-access
- **Evaluate algorithms by running them on a particular sequence of memory references (reference string) and computing the number of page faults on that sequence**
  - Reference string is just page numbers, not full addresses
    - Example reference string: 7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1
  - Repeated access to the same page does not cause a page fault
  - The more physical memory a system has, the more frames it has ... resulting in fewer page-faults

# Page Replacement Algorithms

---

- **Many different page replacement algorithms exist**
  - FIFO Page Replacement
  - Optimal Page Replacement (theoretical best case)
  - Least-Recently Used (LRU) Page Replacement
  - Counting-Based Page Replacement

# FIFO Page Replacement

---

- **A very simple page replacement algorithm**

- When a page is brought into physical memory, insert a reference to that page into a FIFO
- When a page must be replaced, replace the oldest page first (i.e. the page at the head of the FIFO)

- **Pros:**

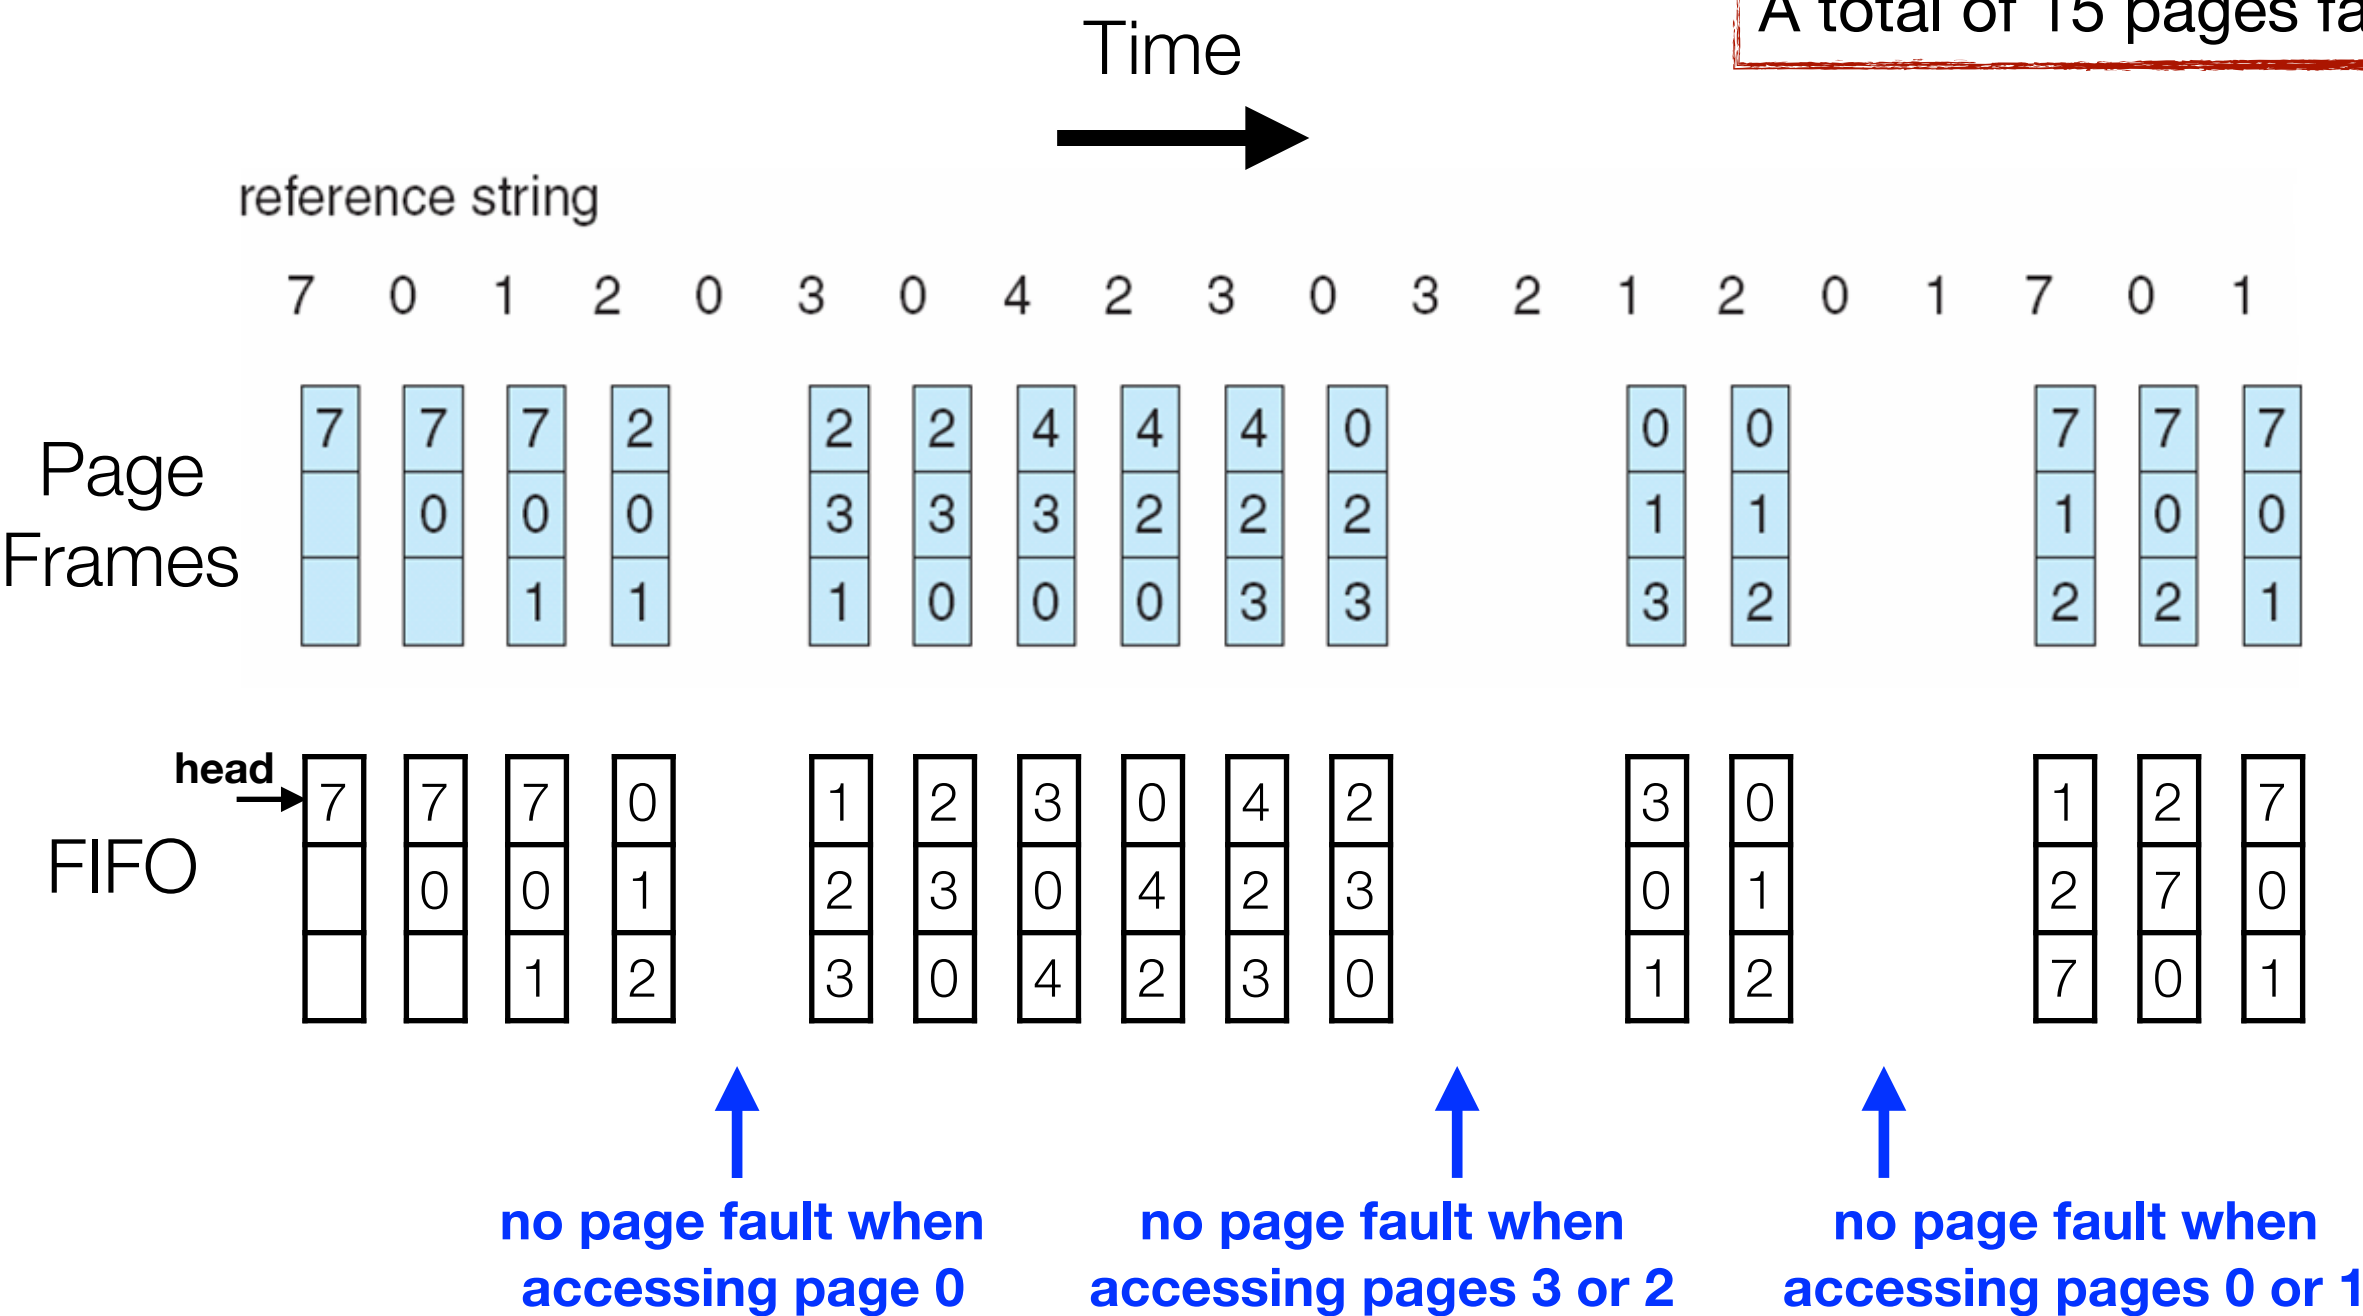
- Very easy to understand and program

- **Cons:**

- Performance may not be very good
- May result in a high number of page faults

# FIFO Page Replacement Example

A total of 15 pages faults

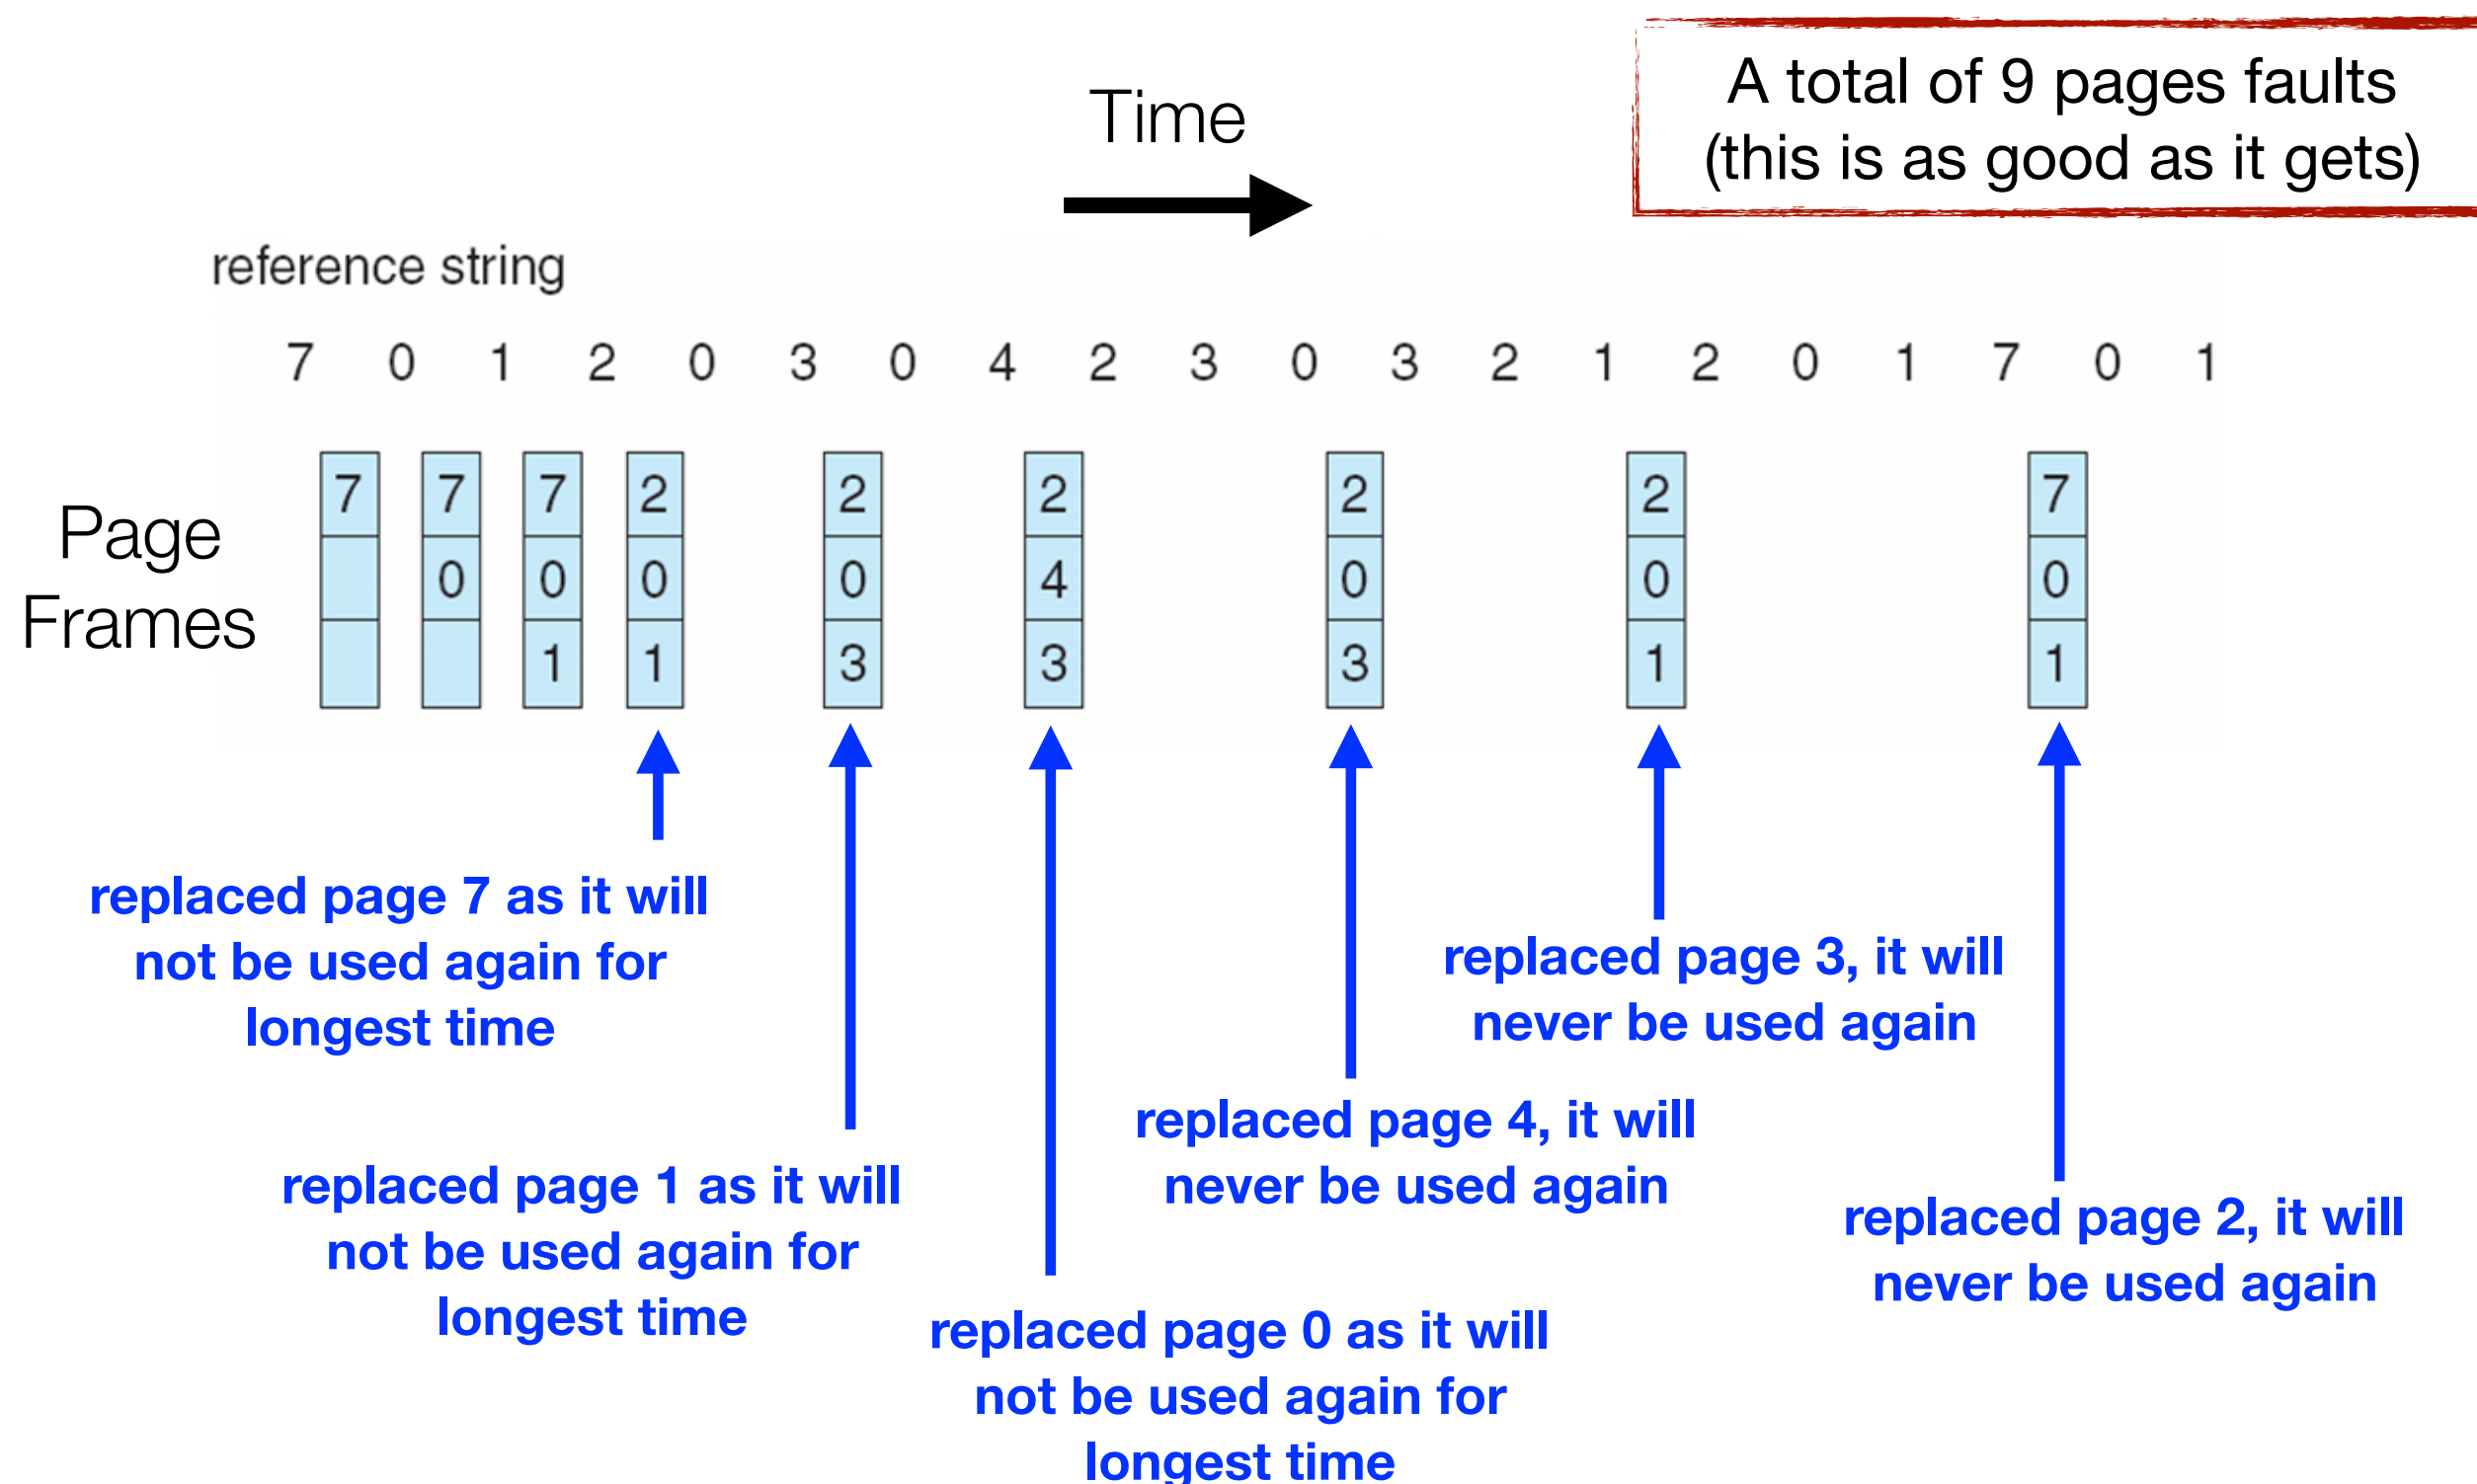


# Optimal Page Replacement

---

- **An optimal page replacement algorithm can be described very simply:**
  - *Replace the page that will not be used for longest period of time*
- **Guarantees the lowest possible page-fault rate for a fixed number of frames**
- **Sadly, it is not possible to know which page won't be used for the longest period of time (can't see the future)**
- **Optimal algorithm is still very useful**
  - Used for measuring how well other algorithm performs
  - How close are other algorithms to optimal?

# Optimal Page Replacement Example



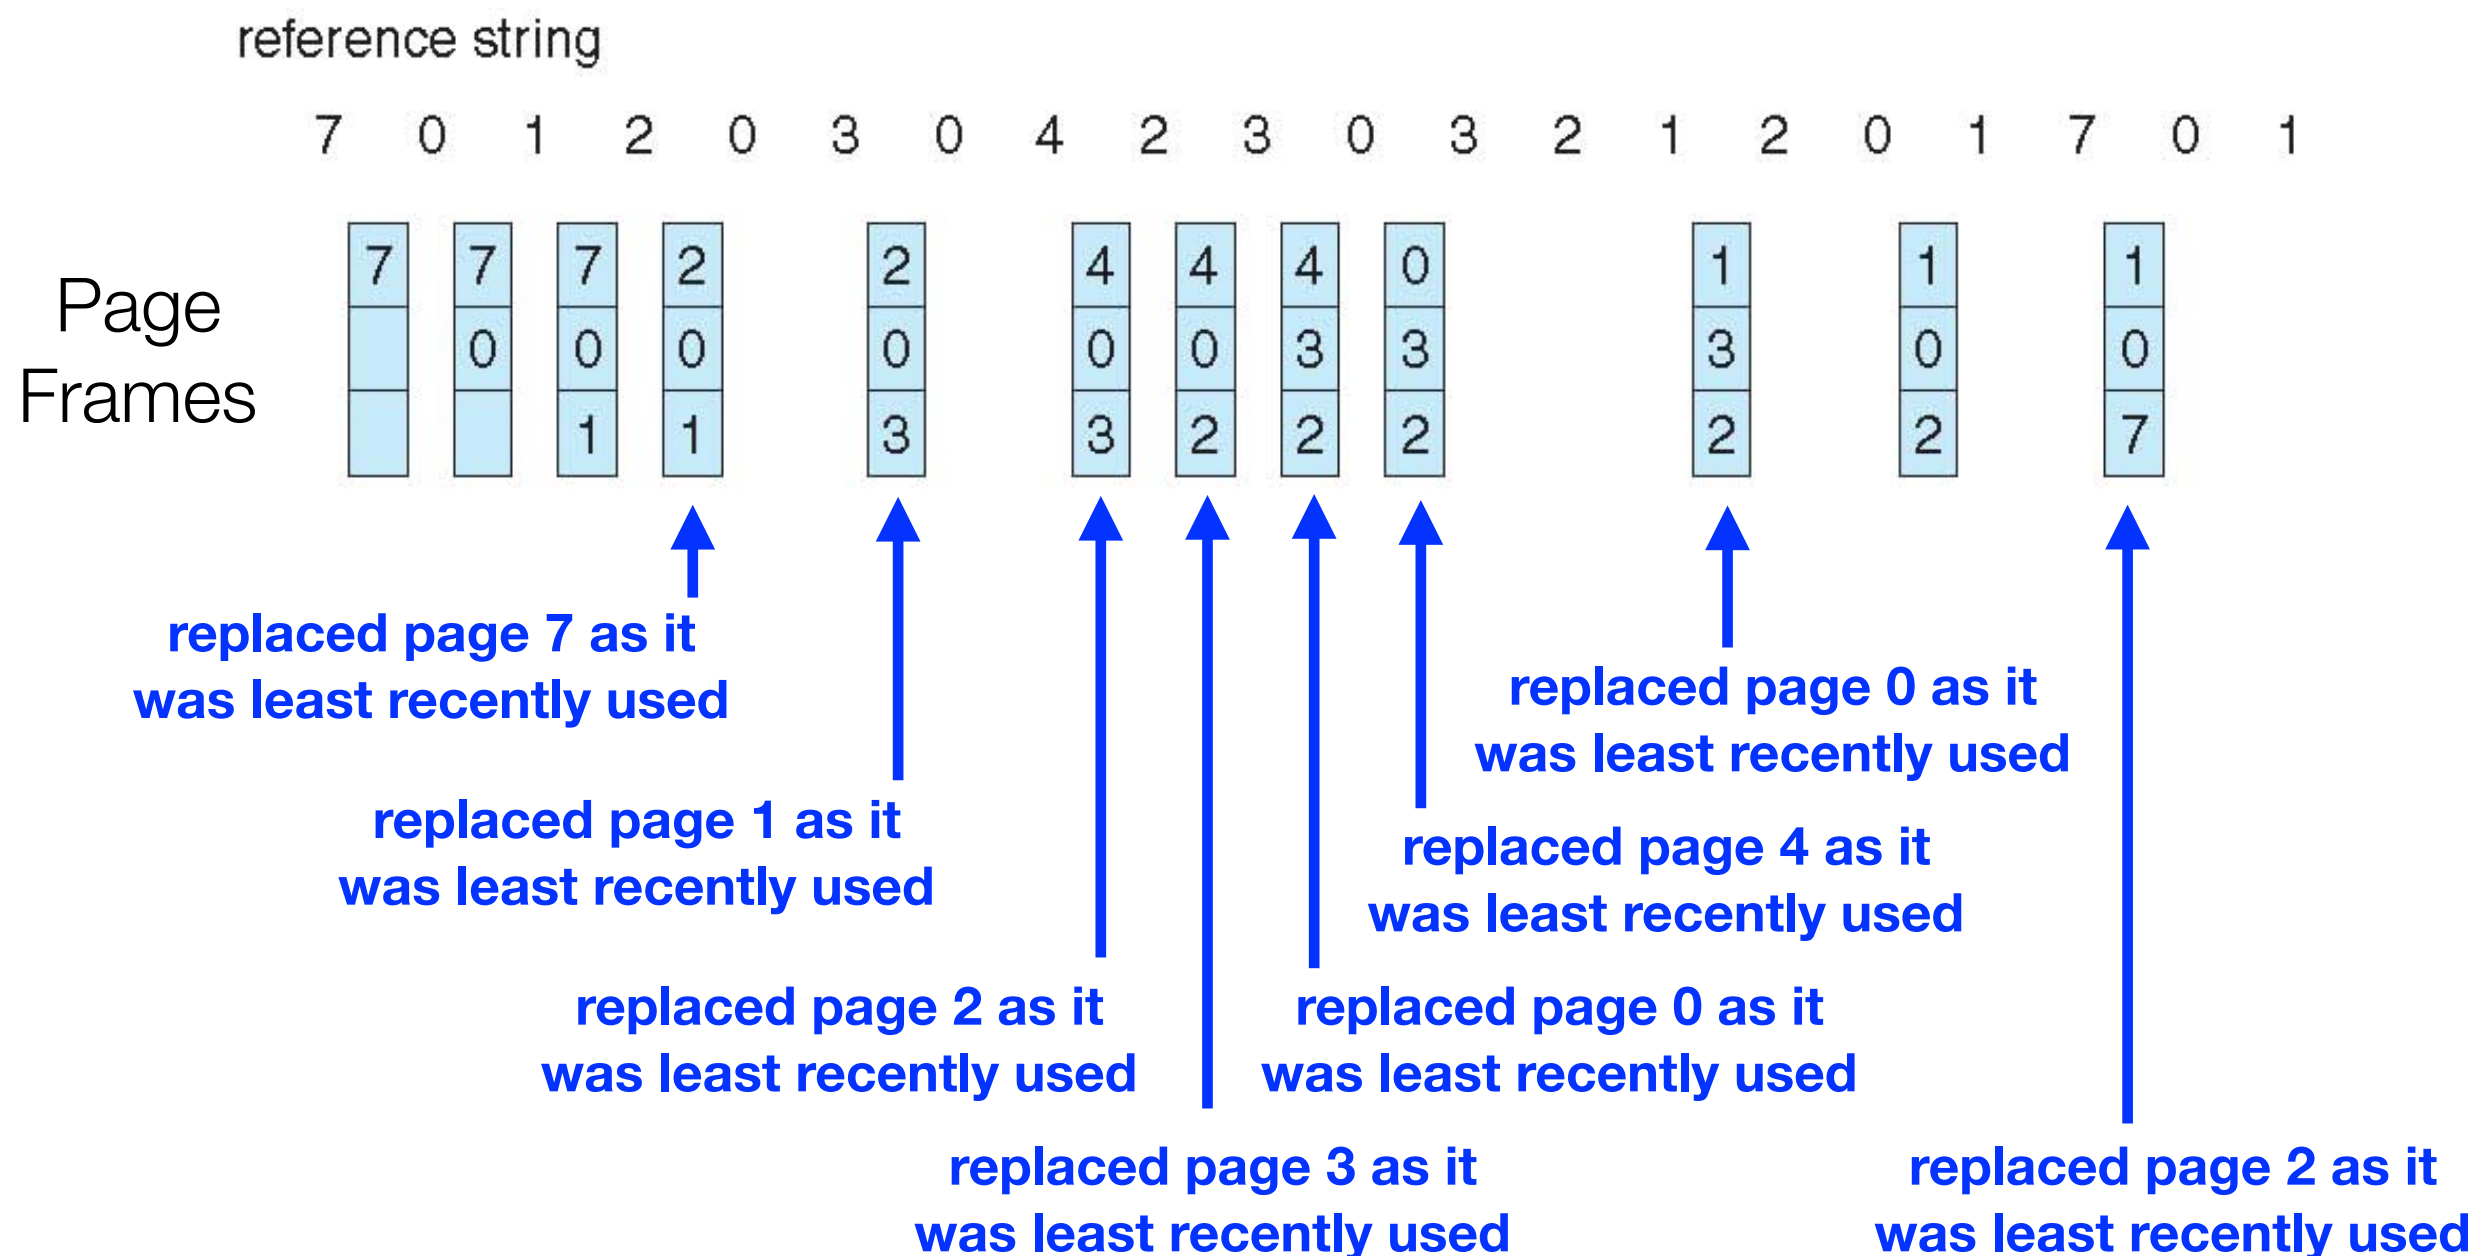
# Least Recently Used (LRU) Algorithm

---

- **Interestingly, the same number of page-faults occur if a reference string is read either forwards or backwards and used to access pages**
  - Because of this property, it is possible to use past knowledge of page use rather than future knowledge
- **Least Recently Used (LRU) Replacement Algorithm:**
  - Replace page that has not been used in the most amount of time (i.e. replace the page that is the least recently used page)
    - Associate time of last use with each page
- **Generally good algorithm and used frequently**

# Least Recently Used (LRU) Algorithm Example

A total of 12 pages faults;  
better than FIFO, but not optimal





# Least Recently Used (LRU) Algorithm

---

- **How should a least-recently-used algorithm be implemented?**

- Option #1 - Counter implementation

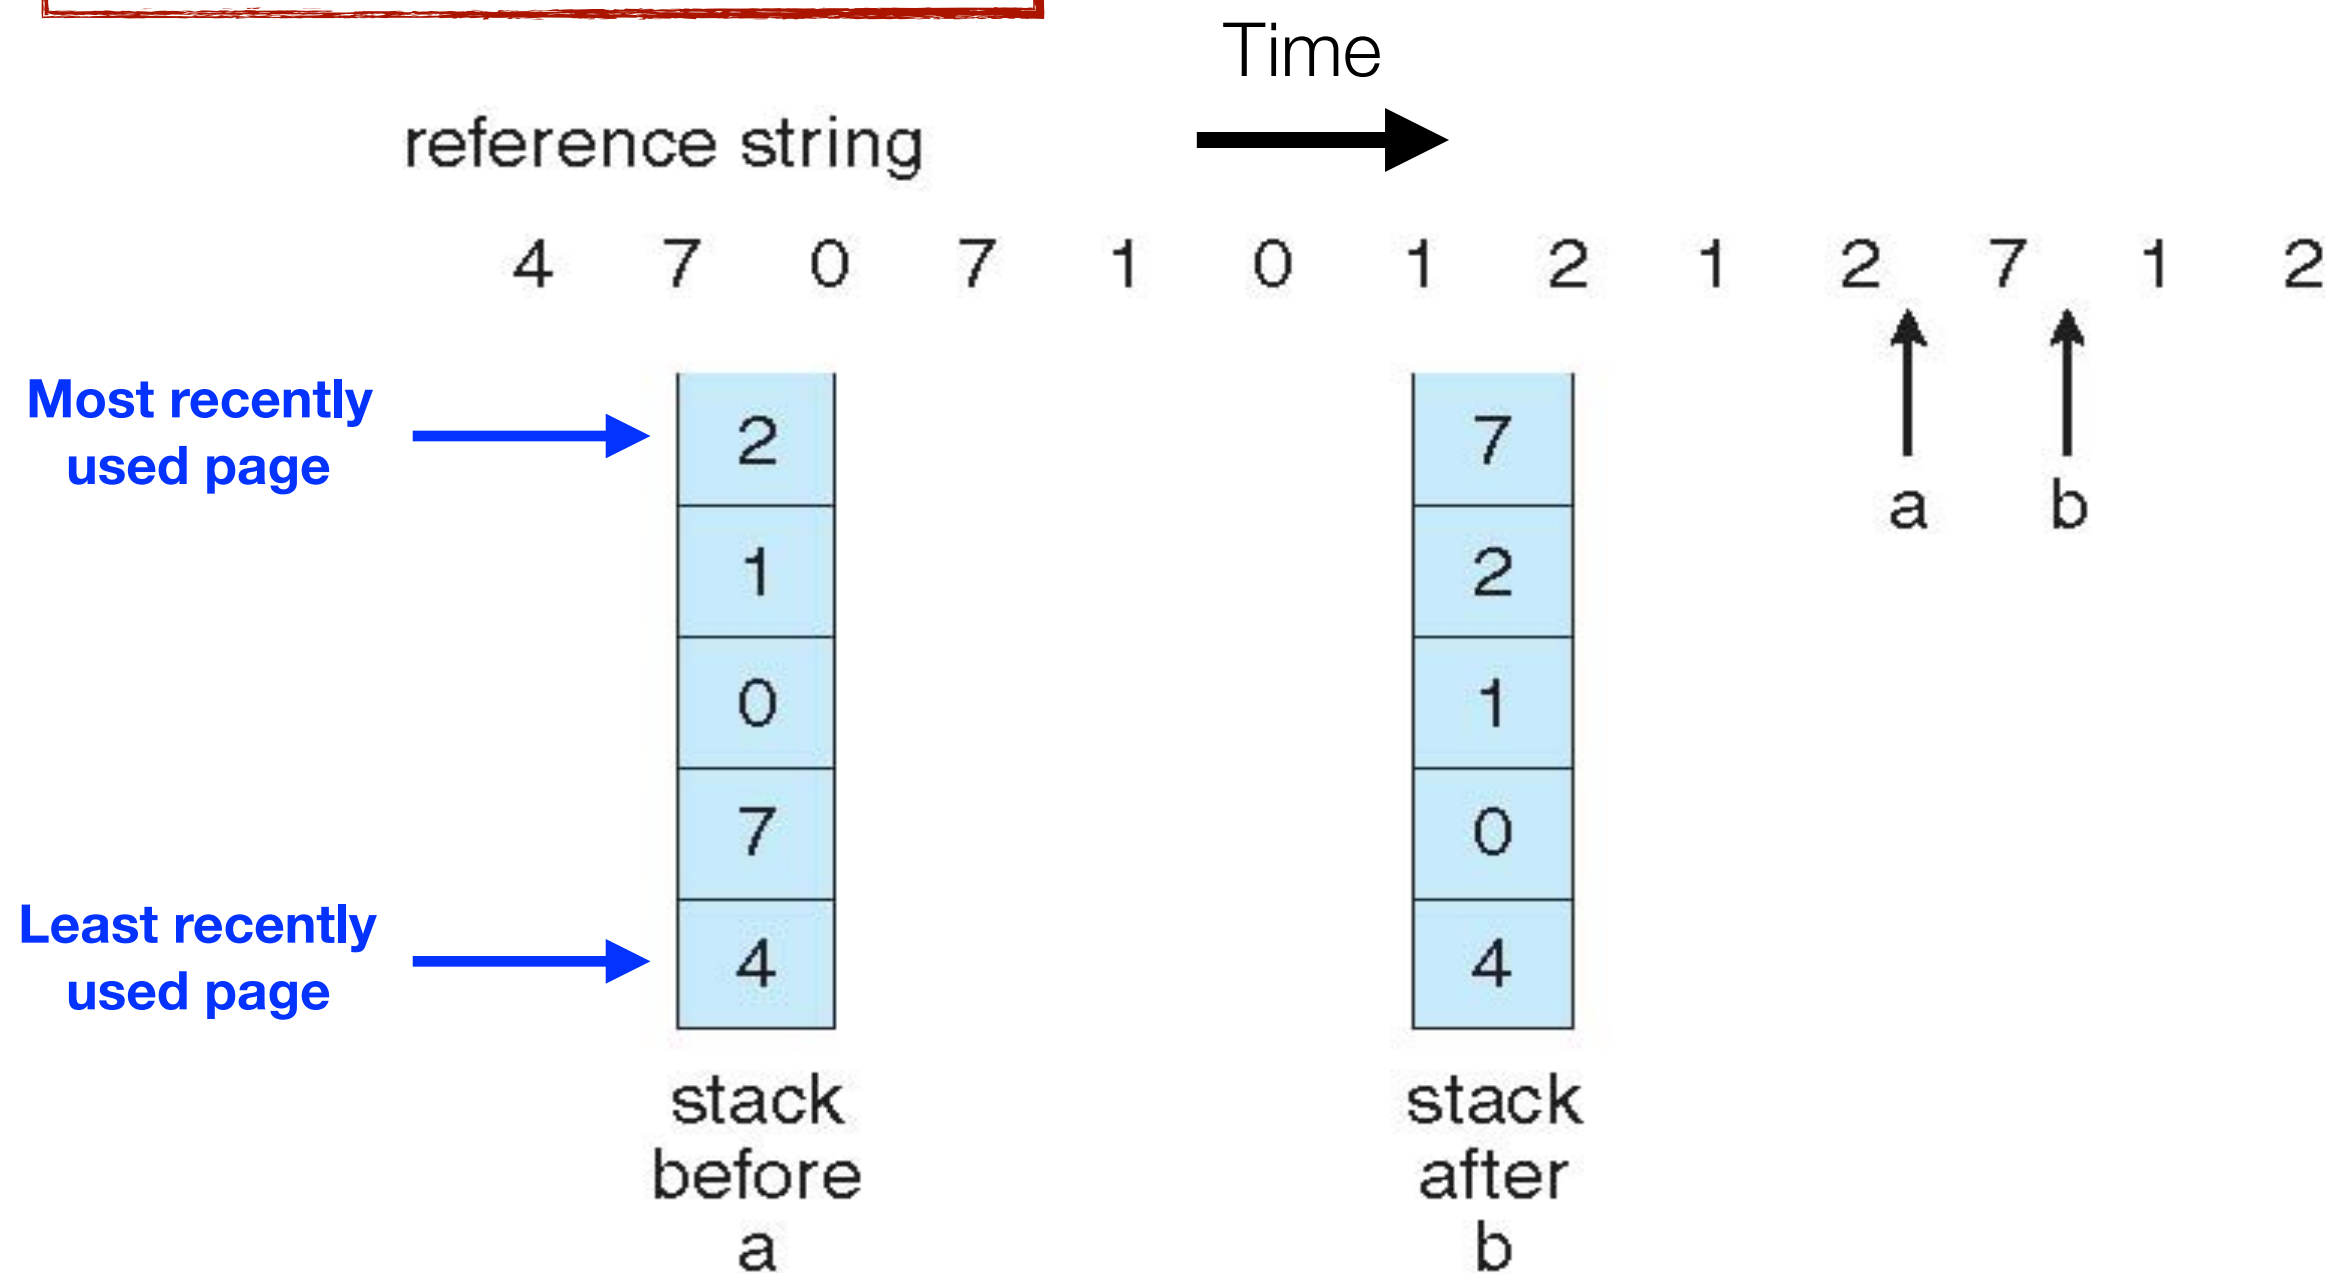
- Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
    - When a page needs to be replaced, look at the counters to find smallest value
      - Requires a potentially lengthy search through the page table (SLOW)

- Option #2 - Stack implementation

- Keep a stack of page numbers in a doubly linked list
    - If a page is referenced, move it to the top of the stack (not a pure stack implementation)
      - Most recently used page is always at the top of the stack; the least recently used page is always at the bottom of the stack
    - Finding and moving page numbers in the stack takes time (updates to the stack must be done on every memory reference) (SLOW)

# Stack Implementation

Using a Stack to Record the Most Recent Page References



# LRU Approximation Algorithms

---

- **The LRU algorithm as described previous is too slow, even with specialized hardware**
- **Instead of finding exactly which page was least recently used, an approximation will do**
- **One possible approach for approximating the least recently used page is to use a reference bit**
  - Associate a reference bit with each page, initially = 0
  - When page is referenced, set bit to 1
  - When it is time to replace a page, replace any that have reference bit = 0 (if one exists)
  - Can also use multiple reference bits to maintain a longer history and thus more closely approximate a true LRU algorithm

# Counting-Based Page Replacement

---

- **Keep a counter of the number of references that have been made to each page**
  - Not common
- **Least-frequently-used (LFU) algorithm:** replaces the page with smallest count
- **Most-frequently-used (MFU) algorithm:** replaces the page with the largest count
  - Based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Speeding Up Page Replacement

---

- **Always maintain a pool of free frames**
  - No need to search for a free frame when the page-fault occurs
  - Load page into free frame and select victim to evict and add to free pool
  - When convenient, evict victim
    - No longer have to wait for victim frame to get paged out before new page can get paged in
- **Possibly, keep list of modified pages**
  - When backing store is otherwise idle, write pages and set to non-dirty
- **Possibly, keep free frame contents intact and note what is in them**
  - If referenced again before reused, no need to load contents again from disk
  - Generally useful to reduce penalty if wrong victim frame selected

# Allocation of Frames

---

- **Each process needs some minimum number of frames that it cannot run without**
  - Even a single instruction may require more than a single frame of physical memory
- **Example: IBM 370 – 6 pages to handle SS MOVE instruction:**
  - Instruction is 6 bytes, and therefore might span 2 pages
  - 2 pages to handle for *from*
  - 2 pages to handle for *to*
- **Maximum number of frames that can be allocated is the total number of frames in the system**
- **Two major allocation frame schemes**
  - Fixed allocation
    - Equal allocation
    - Proportional allocation
  - Priority allocation

# Types of Fixed Allocation

---

- **Equal allocation** – every process in the system is allocated an equal share of the available frames
  - If the number of frames is  $m$  and the number of processes is  $n$ , each process is allocated  $m/n$  frames
  - For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process  $100/5 = 20$  frames
    - Maybe keep some as free frame buffer pool for faster paging
  - Pros: Easy to implement
  - Cons:
    - Why allocate 20 frames to a process that might actually only need 5?
    - Can be very wasteful
    - A higher priority process doesn't get any more frames than a lower priority process

# Types of Fixed Allocation (Cont.)

---

- **Proportional allocation** – Allocate available frames according to the size of a process

- Larger processes are allocated more frames than smaller processes
- Dynamic as degree of multiprogramming change

—  $s_i$  = size of process  $p_i$

—  $S = \sum s_i$

—  $m$  = total number of frames

—  $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$

- Pros: Better allocation scheme than equal allocation
- Cons: A higher priority process doesn't get any more frames than a lower priority process



# Priority Allocation

---

- **Use a proportional allocation scheme using process priorities rather than process size**
- **If process  $P_i$  generates a page fault,**
  - Select for replacement one of its frames (local replacement)
  - Select for replacement a frame from a process with lower priority number (global replacement)

# Global vs. Local Allocation

---

- **Global replacement** – process selects a replacement frame from the set of *ALL* frames; one process can take a frame from another
  - Process execution time can vary greatly due to another process stealing frames
  - Tends to have greater throughput (the number of processes that complete execution per unit time)
  - More common than local replacement
- **Local replacement** – each process selects a replacement frame from only its own set of allocated frames
  - Per-process performance is more consistent
  - But possibly underutilized memory


# Thrashing

---

- **If a process does not have “enough” pages, the page-fault rate is very high**
  - Page fault to get page
  - Replace existing frame
  - But quickly need replaced frame back
  - This leads to:
    - Low CPU utilization
    - Operating system thinking that it needs to increase the degree of multiprogramming
    - Another process added to the system
- **Thrashing  $\equiv$  a process is busy swapping pages in and out (spends more time paging than executing)**

## Thrashing (Cont.)

---



# Demand Paging and Thrashing

---

- **Why does demand paging work?**  
**Locality model**
  - Process migrates from one locality to another
  - Localities may overlap
- **Why does thrashing occur?**  
 **$\Sigma$  size of locality > total memory size**
  - Limit effects by using local or priority page replacement

# Locality In A Memory-Reference Pattern

